

```
import gym

import numpy as np

import tensorflow as tf

from tensorflow.keras import layers


# Define the Checkers Environment

class CheckersEnv(gym.Env):

    def __init__(self):

        # Define your environment here

        pass


    def step(self, action):

        # Execute one time step within the environment

        pass


    def reset(self):

        # Reset the state of the environment to an initial state

        pass


    def render(self, mode='human'):

        # Render the environment to the screen

        pass


# Define the RL Agent

class DQNAgent:

    def __init__(self, state_size, action_size):

        self.state_size = state_size

        self.action_size = action_size

        self.memory = [] # Experience Replay Memory

        self.gamma = 0.95 # Discount rate

        self.epsilon = 1.0 # Exploration rate

        self.epsilon_min = 0.01
```

```
self.epsilon_decay = 0.995
```

```
self.model = self._build_model()
```

```
def _build_model(self):
```

```
    model = tf.keras.Sequential([
```

```
        layers.Dense(24, input_shape=(self.state_size,), activation='relu'),
```

```
        layers.Dense(24, activation='relu'),
```

```
        layers.Dense(self.action_size, activation='linear')
```

```
    ])
```

```
    model.compile(loss='mse', optimizer=tf.keras.optimizers.Adam(lr=0.001))
```

```
    return model
```

```
def remember(self, state, action, reward, next_state, done):
```

```
    self.memory.append((state, action, reward, next_state, done))
```

```
def act(self, state):
```

```
    if np.random.rand() <= self.epsilon:
```

```
        return np.random.choice(self.action_size)
```

```
    act_values = self.model.predict(state)
```

```
    return np.argmax(act_values[0])
```

```
def replay(self, batch_size):
```

```
    minibatch = random.sample(self.memory, batch_size)
```

```
    for state, action, reward, next_state, done in minibatch:
```

```
        target = reward
```

```
        if not done:
```

```
            target = (reward + self.gamma * np.amax(self.model.predict(next_state)[0]))
```

```
        target_f = self.model.predict(state)
```

```
        target_f[0][action] = target
```

```
        self.model.fit(state, target_f, epochs=1, verbose=0)
```

```
    if self.epsilon > self.epsilon_min:
```

```
        self.epsilon *= self.epsilon_decay
```

```

# Initialize the environment and agent

env = CheckersEnv()

state_size = env.observation_space.shape[0]

action_size = env.action_space.n

agent = DQNAgent(state_size, action_size)


# Train the agent

EPISODES = 1000

for e in range(EPISODES):

    state = env.reset()

    state = np.reshape(state, [1, state_size])

    for time in range(500):

        # env.render() # Uncomment if you want to render the environment

        action = agent.act(state)

        next_state, reward, done, _ = env.step(action)

        reward = reward if not done else -10

        next_state = np.reshape(next_state, [1, state_size])

        agent.remember(state, action, reward, next_state, done)

        state = next_state

    if done:

        print("episode: {}/{}", score: {}, e: {:.2}"

              .format(e, EPISODES, time, agent.epsilon))

        break

if len(agent.memory) > batch_size:

    agent.replay(batch_size)

```